

Reg.No.: 210701263

Exp. No.: 3

Map Reduce program to process a weather dataset

AIM:

To implement MapReduce program to process a weather dataset.

Procedure:

Step 1: Create Data File:

Create a file named "word_count_data.txt" and populate it with text data that you wish to analyse.
Login with your hadoop user.

Download the dataset (weather data)

Output:

```
subbu@subbu: ~/exp3
GNU nano 6.2 dataset.txt
23907 20150101 2.423 -98.08 30.62 2.2 -0.6 0.8 0.9 7.0 >
23907 20150102 2.423 -98.08 30.62 3.5 1.3 2.4 2.2 10.2 >
23907 20150103 2.423 -98.08 30.62 15.9 2.3 9.1 7.5 3.1 >
23907 20150104 2.423 -98.08 30.62 9.2 -1.3 3.9 4.2 0.0 >
23907 20150105 2.423 -98.08 30.62 10.9 -3.7 3.6 2.6 0.0 >
23907 20150106 2.423 -98.08 30.62 20.2 2.9 11.6 10.9 0.0 >
23907 20150107 2.423 -98.08 30.62 10.9 -3.4 3.8 4.5 0.0 >
23907 20150108 2.423 -98.08 30.62 0.6 -7.9 -3.6 -3.3 0.0 >
23907 20150109 2.423 -98.08 30.62 2.0 0.1 1.0 0.8 0.0 >
23907 20150110 2.423 -98.08 30.62 0.5 -2.0 -0.8 -0.6 3.9 >
23907 20150111 2.423 -98.08 30.62 10.9 0.0 5.4 4.4 2.6 >
23907 20150112 2.423 -98.08 30.62 6.5 1.4 4.0 4.3 0.0 >
23907 20150113 2.423 -98.08 30.62 3.0 -0.7 1.1 1.2 0.0 >
23907 20150114 2.423 -98.08 30.62 2.9 0.9 1.9 1.8 0.7 >
23907 20150115 2.423 -98.08 30.62 13.2 1.2 7.2 6.4 0.0 >
23907 20150116 2.423 -98.08 30.62 16.7 3.5 10.1 9.9 0.0 >
23907 20150117 2.423 -98.08 30.62 19.5 5.0 12.2 12.3 0.0 >
23907 20150118 2.423 -98.08 30.62 20.9 7.6 14.3 13.7 0.0 >
23907 20150119 2.423 -98.08 30.62 23.9 6.7 15.3 14.3 0.0 >
23907 20150120 2.423 -98.08 30.62 26.0 9.5 17.8 15.9 0.0 >
```

Step 2: Mapper Logic - mapper.py:

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

Reg.No.: 210701263

```
nano mapper.py
```

```
# Copy and paste the mapper.py code
```

```
#!/usr/bin/env python
```

```
import sys
```

```
# input comes from STDIN (standard input)
```

```
# the mapper will get daily max temperature and group it by month. so output will be  
(month,daily_max_temperature)
```

```
for line in sys.stdin:
```

```
    # remove leading and trailing whitespace
```

```
    line = line.strip()    # split
```

```
the line into words    words =
```

```
line.split()
```

```
    #See the README hosted on the weather website which help us understand how each  
position represents a column    month = line[10:12]    daily_max = line[38:45]    daily_max  
= daily_max.strip()
```

```
    # increase counters    for  
word in words:
```

```
    # write the results to STDOUT (standard output);
```

```
    # what we output here will be go through the shuffle process and then
```

```
    # be the input for the Reduce step, i.e. the input for reducer.py
```

```
    #
```

```
    # tab-delimited; month and daily max temperature as output
```

```
print ('%s\t%s' % (month ,daily_max))
```

```
.
```

Step 3: Reducer Logic - reducer.py:

Create a file named "reducer.py" to implement the logic for the reducer. The reducer will aggregate the occurrences of each word and generate the final output.

```
nano reducer.py
```

```
# Copy and paste the reducer.py code
```

```
reducer.py
```

```
#!/usr/bin/env python
```

```
from operator import itemgetter import sys
```

```
#reducer will get the input from stdid which will be a collection of key, value(Key=month , value=  
daily_max_temperature)
```

```
#reducer logic: will get all the daily max temperature for a month and find max temperature for the  
month
```

```
#shuffle will ensure that key are sorted(month)
```

Reg.No.: 210701263

```
current_month = None
current_max = 0
month = None

# input comes from STDIN for
line in sys.stdin:
    # remove leading and trailing whitespace    line
    = line.strip()
    # parse the input we got from mapper.py    month,
    daily_max = line.split('\t', 1)

    # convert daily_max (currently a string) to float    try:
        daily_max = float(daily_max)    except
ValueError:
    # daily_max was not a number, so silently
    # ignore/discard this line
    continue

    # this IF-switch only works because Hadoop shuffle process sorts map output
    # by key (here: month) before it is passed to the reducer
    if current_month == month:        if daily_max > current_max:
        current_max = daily_max    else:        if current_month:
            # write result to STDOUT
            print ('%s\t%s' % (current_month, current_max))
        current_max = daily_max
        current_month = month

# output of the last month if current_month == month:
print ('%s\t%s' % (current_month, current_max))
```

Step 4: Prepare Hadoop Environment:

Start the Hadoop daemons and create a directory in HDFS to store your data.

```
start-all.sh
```

Step 6: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.

```
chmod 777 mapper.py reducer.py
```

Reg.No.: 210701263

Step 7: Run the program using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access.

Then run the program using Hadoop Streaming.

```
hadoop fs -mkdir -p /weatherdata
```

```
hadoop fs -copyFromLocal /home/sx/Downloads/dataset.txt /weatherdata
```

```
hdfs dfs -ls /weatherdata
```

```
hadoop jar /home/sx/hadoop-3.2.3/share/hadoop/tools/lib/hadoop-streaming-3.2.3.jar \  
-input /weatherdata/dataset.txt \  
-output /weatherdata/output \  
-file "/home/sx/Downloads/mapper.py" \  
-mapper "python3 mapper.py" \  
-file "/home/sx/Downloads/reducer.py" \  
-reducer "python3 reducer.py"
```

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/outputfile.txt
```

Step 8: Check Output:

Check the output of the program in the specified HDFS output directory.

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/output/ /part-00000
```

Reg.No.: 210701263

```
subbu@subbu: ~/exp3
demo_pig.pig      pig_1726811332057.log  udf_example.pig
pig_1726806667729.log  sample.txt      uppercase_udf.py
subbu@subbu:~/exp4$ cd ..
subbu@subbu:~$ cd exp3
subbu@subbu:~/exp3$ ls
cmd.txt  dataset.txt  hadoop-streaming-3.3.6.jar  mapper.py  reducer.py
subbu@subbu:~/exp3$ dataset.txt
dataset.txt: command not found
subbu@subbu:~/exp3$ nano dataset.txt
subbu@subbu:~/exp3$ chmod 777 mapper.py reducer.py
subbu@subbu:~/exp3$ hdfs dfs -cat /exp3/output1/part-00000
01      26.5
02      26.6
03      29.1
04      30.8
05      31.1
06      33.6
07      38.5
08      40.2
09      36.5
10      36.9
11      27.6
12      25.9
subbu@subbu:~/exp3$
```

After copy and paste the above output in your local file give the below command to remove the directory from hdfs : `hadoop fs -rm -r /weatherdata/output`

Result:

Thus, the program for weather dataset using Map Reduce has been executed successfully.